

In [35]:

```
import pandas as pd
import numpy as np
import statsmodels.api as sm
import sklearn.metrics as metrics
from sklearn.utils import resample
from sklearn.linear_model import LinearRegression
from sklearn.model_selection import train_test_split
```

Question 1.

Estimate the MSE of the model using the traditional approach. That is, fit the linear regression model using the entire dataset and calculate the mean squared error for the entire dataset. Present and discuss your results at a single, high level.

In [2]:

```
biden = pd.read_csv('nes2008.csv')
biden.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1807 entries, 0 to 1806
Data columns (total 6 columns):
biden      1807 non-null int64
female     1807 non-null int64
age        1807 non-null int64
educ       1807 non-null int64
dem        1807 non-null int64
rep        1807 non-null int64
dtypes: int64(6)
memory usage: 84.8 KB
```

In [6]:

```
x = biden[['female', 'age', 'educ', 'dem', 'rep']]
y = biden['biden']

x = sm.add_constant(x)
modell = sm.OLS(y, x).fit()
modell.summary()
```

```
/opt/anaconda3/lib/python3.7/site-packages/numpy/core/fromnumeric.py:2495: FutureWarning: Method .ptp is deprecated and will be removed in a future version. Use numpy.ptp instead.
```

```
    return ptp(axis=axis, out=out, **kwargs)
```

Out[6]:

OLS Regression Results

Dep. Variable:	biden	R-squared:	0.282
Model:	OLS	Adj. R-squared:	0.280
Method:	Least Squares	F-statistic:	141.1
Date:	Mon, 03 Feb 2020	Prob (F-statistic):	1.50e-126
Time:	13:51:55	Log-Likelihood:	-7966.6
No. Observations:	1807	AIC:	1.595e+04
Df Residuals:	1801	BIC:	1.598e+04
Df Model:	5		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
const	58.8113	3.124	18.823	0.000	52.683	64.939
female	4.1032	0.948	4.327	0.000	2.243	5.963
age	0.0483	0.028	1.708	0.088	-0.007	0.104
educ	-0.3453	0.195	-1.773	0.076	-0.727	0.037
dem	15.4243	1.068	14.442	0.000	13.330	17.519
rep	-15.8495	1.311	-12.086	0.000	-18.421	-13.278

Omnibus:	87.979	Durbin-Watson:	1.996
Prob(Omnibus):	0.000	Jarque-Bera (JB):	101.940
Skew:	-0.533	Prob(JB):	7.31e-23
Kurtosis:	3.466	Cond. No.	348.

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

In [8]:

```
y_pred = model1.predict(x)
metrics.mean_squared_error(y, y_pred)
```

Out[8]:

395.2701692786484

The mean squared error for the entire dataset is 395.27. In the linear regression result, r-squared is 0.282, representing that this model only explains 28.2% variation in the dataset. For the features, the p-values of female, dem, and rep are pretty small. Therefore, it's fair to say that we have 99% confidence level to say that female and demographic have more active attitudes towards Joe Biden than male and republican. The coefficient of age is positive, education is negative. But the p-values are not that small.

Question 2.

In [12]:

```
x_train, x_test, y_train, y_test = train_test_split(x, y, random_state = 0, test_size=0.2)
```

In [13]:

```
x_train = sm.add_constant(x_train)
model2 = sm.OLS(y_train, x_train).fit()
model2.summary()
```

```
/opt/anaconda3/lib/python3.7/site-packages/numpy/core/fromnumeric.py:2
495: FutureWarning: Method .ptp is deprecated and will be removed in a
future version. Use numpy.ptp instead.
```

```
    return ptp(axis=axis, out=out, **kwargs)
```

Out[13]:

OLS Regression Results

Dep. Variable:	biden	R-squared:	0.290
Model:	OLS	Adj. R-squared:	0.286
Method:	Least Squares	F-statistic:	73.22
Date:	Mon, 03 Feb 2020	Prob (F-statistic):	2.42e-64
Time:	13:53:58	Log-Likelihood:	-3998.3
No. Observations:	903	AIC:	8009.
Df Residuals:	897	BIC:	8037.
Df Model:	5		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
const	61.1071	4.479	13.642	0.000	52.316	69.899
female	1.6538	1.368	1.209	0.227	-1.032	4.339
age	0.0362	0.041	0.882	0.378	-0.044	0.117
educ	-0.3524	0.278	-1.268	0.205	-0.898	0.193
dem	15.7318	1.541	10.207	0.000	12.707	18.757
rep	-17.8384	1.896	-9.407	0.000	-21.560	-14.117

Omnibus:	21.887	Durbin-Watson:	2.026
Prob(Omnibus):	0.000	Jarque-Bera (JB):	22.939
Skew:	-0.389	Prob(JB):	1.04e-05
Kurtosis:	3.069	Cond. No.	346.

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

In [25]:

```
x_test = sm.add_constant(x_test)
y_pred = model2.predict(x_test)
metrics.mean_squared_error(y_test, y_pred)
```

```
/opt/anaconda3/lib/python3.7/site-packages/numpy/core/fromnumeric.py:2
495: FutureWarning: Method .ptp is deprecated and will be removed in a
future version. Use numpy.ptp instead.
    return ptp(axis=axis, out=out, **kwargs)
```

Out[25]:

384.3787194770427

The MSE of question 2 is 384.38, which is different from what we calculate in question 1, 395.27. The dataset that we use to calculate MSE in question 1 is exactly what we use to train the model. However, in question 2, the dataset that we use to calculate MSE is the model has never seen before. Therefore, it is fair that the MSE in question 2 is smaller than question 1. Also, it's possible that the MSE in question 2 is bigger/ equal that question 1, because this depends on how we split the training set and the test set.

Question 3

In [22]:

```
mse1000 = []
for i in range(1000):
    xtra, xte, ytra, yte = train_test_split(x, y, random_state=i, test_size=0.5)
    xtra = sm.add_constant(xtra)
    model = sm.OLS(ytra, xtra).fit()
    #mlresult = mlmodel.fit()
    xte = sm.add_constant(xte)
    y_pred = model.predict(xte)
    mse1000.append(metrics.mean_squared_error(yte, y_pred))
```

```
/opt/anaconda3/lib/python3.7/site-packages/numpy/core/fromnumeric.py:2
495: FutureWarning: Method .ptp is deprecated and will be removed in a
future version. Use numpy.ptp instead.
    return ptp(axis=axis, out=out, **kwargs)
/opt/anaconda3/lib/python3.7/site-packages/numpy/core/fromnumeric.py:2
495: FutureWarning: Method .ptp is deprecated and will be removed in a
future version. Use numpy.ptp instead.
    return ptp(axis=axis, out=out, **kwargs)
/opt/anaconda3/lib/python3.7/site-packages/numpy/core/fromnumeric.py:2
495: FutureWarning: Method .ptp is deprecated and will be removed in a
future version. Use numpy.ptp instead.
    return ptp(axis=axis, out=out, **kwargs)
/opt/anaconda3/lib/python3.7/site-packages/numpy/core/fromnumeric.py:2
495: FutureWarning: Method .ptp is deprecated and will be removed in a
future version. Use numpy.ptp instead.
    return ptp(axis=axis, out=out, **kwargs)
/opt/anaconda3/lib/python3.7/site-packages/numpy/core/fromnumeric.py:2
495: FutureWarning: Method .ptp is deprecated and will be removed in a
future version. Use numpy.ptp instead.
```

In [31]:

```
sum(mse1000) / len(mse1000)
```

Out[31]:

399.6615061066417

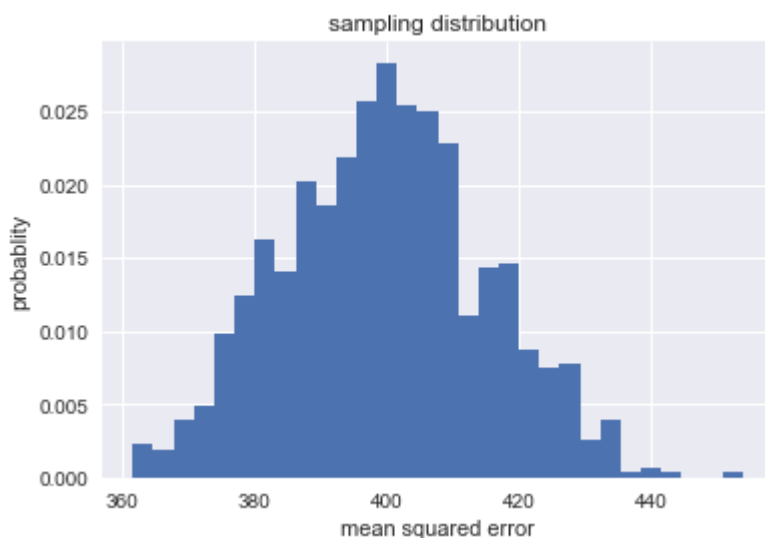
In [33]:

```
import matplotlib.pyplot as plt
%matplotlib inline
plt.hist(mse1000, normed=True, bins=30)
plt.xlabel('mean squared error')
plt.ylabel('probability')
plt.title('sampling distribution')
plt.show()
```

/opt/anaconda3/lib/python3.7/site-packages/ipykernel_launcher.py:3: MatplotlibDeprecationWarning:

The 'normed' kwarg was deprecated in Matplotlib 2.1 and will be removed in 3.1. Use 'density' instead.

This is separate from the ipykernel package so we can avoid doing imports until



The picture looks like a normal distribution, corresponding to the central limit theorem. The mean of MSE is 399.66. Recall that we get 395.27 in question 1 and 384.38 in question 2, 395.27 is closer to 384.38. This makes sense because 395.27 is to fit the entire dataset. And this is a little different from 399.6 because using bootstrap method helps to get close to the population mean.

Question 4

In [40]:

```
boot = resample(biden, replace=True, n_samples=1000, random_state=0)
boot.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 1000 entries, 684 to 1168
Data columns (total 6 columns):
biden      1000 non-null int64
female     1000 non-null int64
age         1000 non-null int64
educ        1000 non-null int64
dem         1000 non-null int64
rep         1000 non-null int64
dtypes: int64(6)
memory usage: 54.7 KB
```

In [41]:

```
x_train = boot[['female', 'age', 'educ', 'dem', 'rep']]
y_train = boot[['biden']]
x_train = sm.add_constant(x_train)
model4 = sm.OLS(y_train, x_train).fit()
model4.summary()
```

```
/opt/anaconda3/lib/python3.7/site-packages/numpy/core/fromnumeric.py:2
495: FutureWarning: Method .ptp is deprecated and will be removed in a
future version. Use numpy.ptp instead.
    return ptp(axis=axis, out=out, **kwargs)
```

Out[41]:

OLS Regression Results

Dep. Variable:	biden	R-squared:	0.291
Model:	OLS	Adj. R-squared:	0.288
Method:	Least Squares	F-statistic:	81.71
Date:	Mon, 03 Feb 2020	Prob (F-statistic):	6.35e-72
Time:	14:10:21	Log-Likelihood:	-4420.4
No. Observations:	1000	AIC:	8853.
Df Residuals:	994	BIC:	8882.
Df Model:	5		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
const	59.5482	4.174	14.266	0.000	51.357	67.739
female	1.8492	1.287	1.437	0.151	-0.676	4.374
age	0.0137	0.038	0.357	0.721	-0.062	0.089
educ	-0.3431	0.261	-1.315	0.189	-0.855	0.169
dem	18.3129	1.448	12.647	0.000	15.471	21.154
rep	-14.3282	1.785	-8.027	0.000	-17.831	-10.825

Omnibus:	31.600	Durbin-Watson:	2.099
Prob(Omnibus):	0.000	Jarque-Bera (JB):	33.771
Skew:	-0.441	Prob(JB):	4.64e-08
Kurtosis:	3.179	Cond. No.	345.

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

In [44]:

```
y_pred = model4.predict(x_train)
metrics.mean_squared_error(y_train, y_pred)
```

Out[44]:

404.57987983796636

The numeric output comparison of question 1 and question 4:

In [45]:

```
model1.summary()
```

Out[45]:

OLS Regression Results

Dep. Variable:	biden	R-squared:	0.282
Model:	OLS	Adj. R-squared:	0.280
Method:	Least Squares	F-statistic:	141.1
Date:	Mon, 03 Feb 2020	Prob (F-statistic):	1.50e-126
Time:	14:11:51	Log-Likelihood:	-7966.6
No. Observations:	1807	AIC:	1.595e+04
Df Residuals:	1801	BIC:	1.598e+04
Df Model:	5		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
const	58.8113	3.124	18.823	0.000	52.683	64.939
female	4.1032	0.948	4.327	0.000	2.243	5.963
age	0.0483	0.028	1.708	0.088	-0.007	0.104
educ	-0.3453	0.195	-1.773	0.076	-0.727	0.037
dem	15.4243	1.068	14.442	0.000	13.330	17.519
rep	-15.8495	1.311	-12.086	0.000	-18.421	-13.278

Omnibus:	87.979	Durbin-Watson:	1.996
Prob(Omnibus):	0.000	Jarque-Bera (JB):	101.940
Skew:	-0.533	Prob(JB):	7.31e-23
Kurtosis:	3.466	Cond. No.	348.

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

In [46]:

```
model4.summary()
```

Out[46]:

OLS Regression Results

Dep. Variable:	biden	R-squared:	0.291
Model:	OLS	Adj. R-squared:	0.288
Method:	Least Squares	F-statistic:	81.71
Date:	Mon, 03 Feb 2020	Prob (F-statistic):	6.35e-72
Time:	14:11:57	Log-Likelihood:	-4420.4
No. Observations:	1000	AIC:	8853.
Df Residuals:	994	BIC:	8882.
Df Model:	5		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
const	59.5482	4.174	14.266	0.000	51.357	67.739
female	1.8492	1.287	1.437	0.151	-0.676	4.374
age	0.0137	0.038	0.357	0.721	-0.062	0.089
educ	-0.3431	0.261	-1.315	0.189	-0.855	0.169
dem	18.3129	1.448	12.647	0.000	15.471	21.154
rep	-14.3282	1.785	-8.027	0.000	-17.831	-10.825

Omnibus:	31.600	Durbin-Watson:	2.099
Prob(Omnibus):	0.000	Jarque-Bera (JB):	33.771
Skew:	-0.441	Prob(JB):	4.64e-08
Kurtosis:	3.179	Cond. No.	345.

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

In general, the numerical outputs of question 1 and question 4 are pretty similar. The R-squared in question 1 is 0.282, and 0.291 in question 4. The difference in coefficients are the bootstrap method has smaller estimates for female and age, and bigger estimates for education, demographic, and republican. A little bit bigger in question 4, but not very much. The p-value of female is 0.151 in question 4, comparing to 0.000 in question 1. The p-value of age gets bigger, you cannot say anything about a 0.721 p-value, and the p-value of education also becomes bigger.

Generally, it's better to use bootstrap method because it's possible that the dataset cannot meet the distribution assumption, so the bootstrap method gives us a more accurate and robust result.

In []: